

Chapter 5.3.1

Image Filtering, Cross-correlation, Convolution

Sources:

- Sonka Textbook
- Gonzalez/Woods DIP textbook

Overview

- Correlation and convolution
- Linear filtering
 - Smoothing, kernels, models
 - Detection
 - Derivatives
- Nonlinear filtering
 - Median filtering
 - Bilateral filtering
 - Neighborhood statistics and nonlocal filtering

Cross Correlation

- Operation on image neighborhood and small ...
 - “mask”, “filter”, “stencil”, “kernel”
- Linear operations within a moving window



Cross Correlation

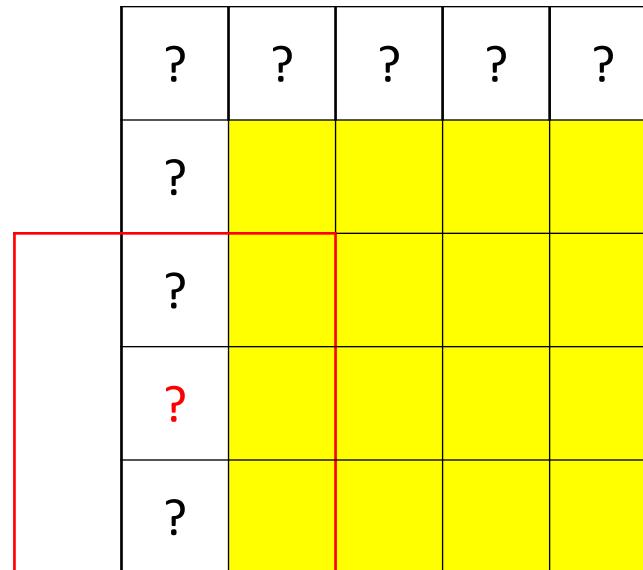
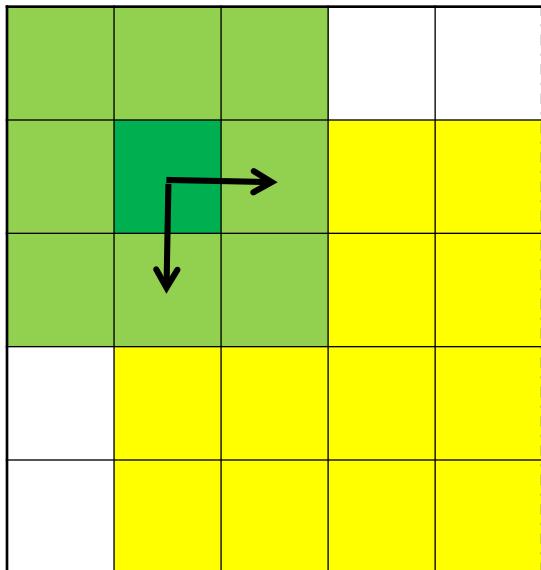
- 1D
$$g(x) = \sum_{s=-a}^a w(s)f(x+s)$$

- 2D
$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t)f(x+s, y+t)$$

$$w(s, t) = \begin{bmatrix} w(-a, -b) & \cdots & & \cdots & w(a, -b) \\ \vdots & & & & \vdots \\ & \cdots & w(0, 0) & \cdots & \\ \vdots & & & & \vdots \\ w(-a, b) & \cdots & & \cdots & w(a, b) \end{bmatrix}$$

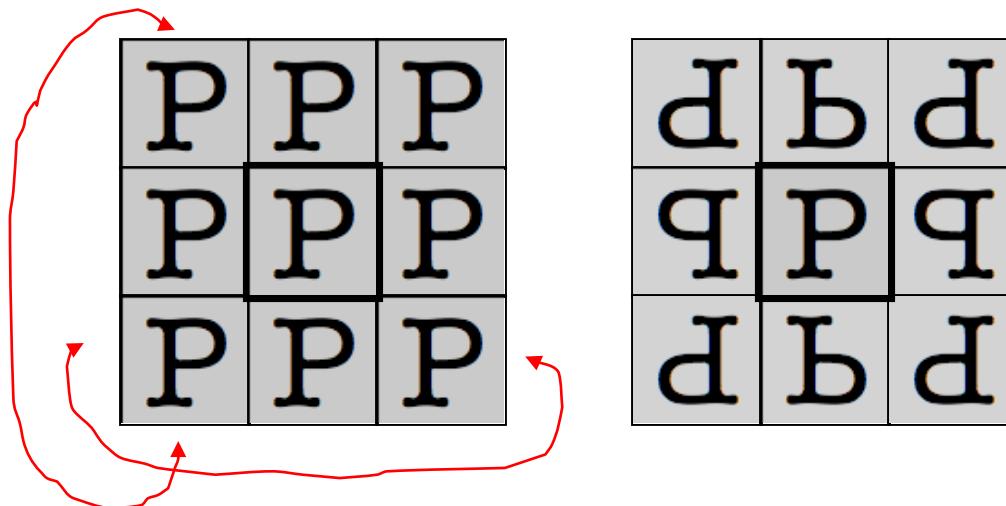
Correlation: Technical Details

- How to filter boundary?



Correlation: Technical Details

- Boundary conditions
 - Boundary not filtered (keep it 0)
 - Pad image with amount (a,b)
 - Constant value or repeat edge values
 - Cyclical boundary conditions
 - Wrap or mirroring



Correlation: Technical Details

- **Boundaries**
 - Can also modify kernel – no longer correlation
- **For analysis**
 - Image domains infinite
 - Data compact (goes to zero far away from origin)

$$g(x, y) = \sum_{s=-\infty}^{\infty} \sum_{t=-\infty}^{\infty} w(s, t) f(x + s, y + t)$$

Correlation: Properties

- Shift invariant

$$g = w \circ f \quad g(x, y) = w(x, y) \circ f(x, y)$$

$$w(x, y) \circ f(x - x_0, y - y_0) = \sum_{s=-\infty}^{\infty} \sum_{t=-\infty}^{\infty} w(s, t) f(x - x_0 + s, y - y_0 + t) = g(x - x_0, y - y_0)$$

Correlation: Properties

- Shift invariant

$$g = w \circ f \quad g(x, y) = w(x, y) \circ f(x, y)$$

$$w(x, y) \circ f(x - x_0, y - y_0) = \sum_{s=-\infty}^{\infty} \sum_{t=-\infty}^{\infty} w(s, t) f(x - x_0 + s, y - y_0 + t) = g(x - x_0, y - y_0)$$

- Linear $w \circ (\alpha e + \beta f) = \alpha w \circ e + \beta w \circ f$

Compact notation

$$C_{wf} = w \circ f$$

Filters: Considerations

- Normalize
 - Sums to one
 - Sums to zero (some cases, see later)
- Symmetry
 - Left, right, up, down
 - Rotational
- Special case: auto correlation

$$C_{ff} = f \circ f$$

Examples 1



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$



$$\begin{matrix} 1 & 1 & 1 \\ 1/9 * & 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Examples 1



$$\begin{matrix} 1 & 1 & 1 \\ 1/9 * & 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Examples 2



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



$$1/25 * \begin{matrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{matrix}$$



Smoothing and Noise

Noisy image



5x5 box filter



Noise Analysis

- Consider an a simple image $I()$ with additive, uncorrelated, zero-mean noise of variance s
- What is the expected rms error of the corrupted image?
- If we process the image with a box filter of size $2a+1$ what is the expected error of the filtered image?

$$\text{RMSE} = \left(\frac{1}{|\mathcal{D}|} \sum_{(x,y) \in \mathcal{D}} \left(\tilde{I}(x,y) - I(x,y) \right)^2 \right)^{\frac{1}{2}}$$

Other Filters

- Disk
 - Circularly symmetric, jagged in discrete case
- Gaussians
 - Circularly symmetric, smooth for large enough stdev
 - Must normalize in order to sum to one
- Derivatives – discrete/finite differences
 - Operators

Gaussian Kernel

```
 $\sigma = 1$ ; Plot[ $\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{x^2}{2\sigma^2}}$ , {x, -4, 4}, ImageSize ->
```

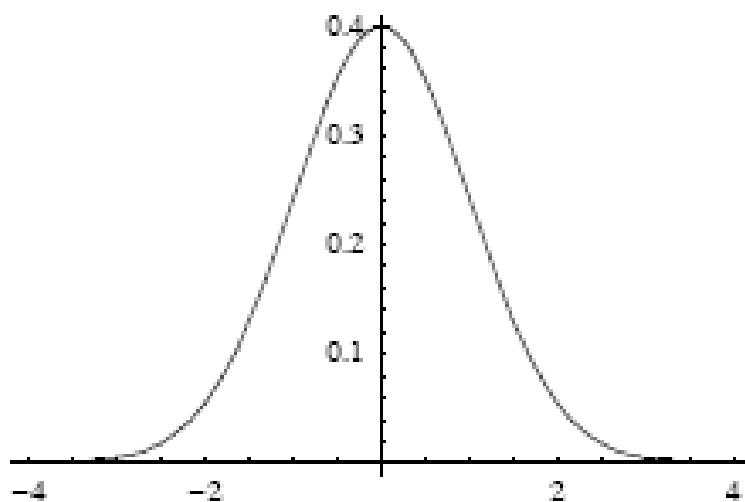


Figure 2.1 The Gaussian kernel with unit standard deviation in 1D.

Gaussian Kernel

$$G_{1D}(x; \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}, \quad G_{2D}(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}, \quad G_{\text{ND}}(\vec{x}; \sigma) = \frac{1}{(\sqrt{2\pi}\sigma)^N} e^{-\frac{|\vec{x}|^2}{2\sigma^2}}$$

Normalization to 1.0

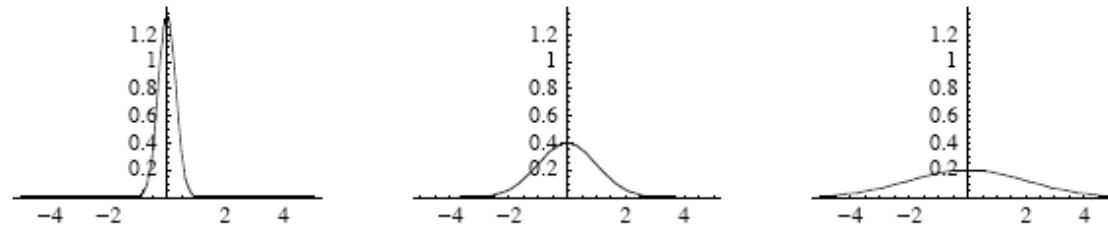
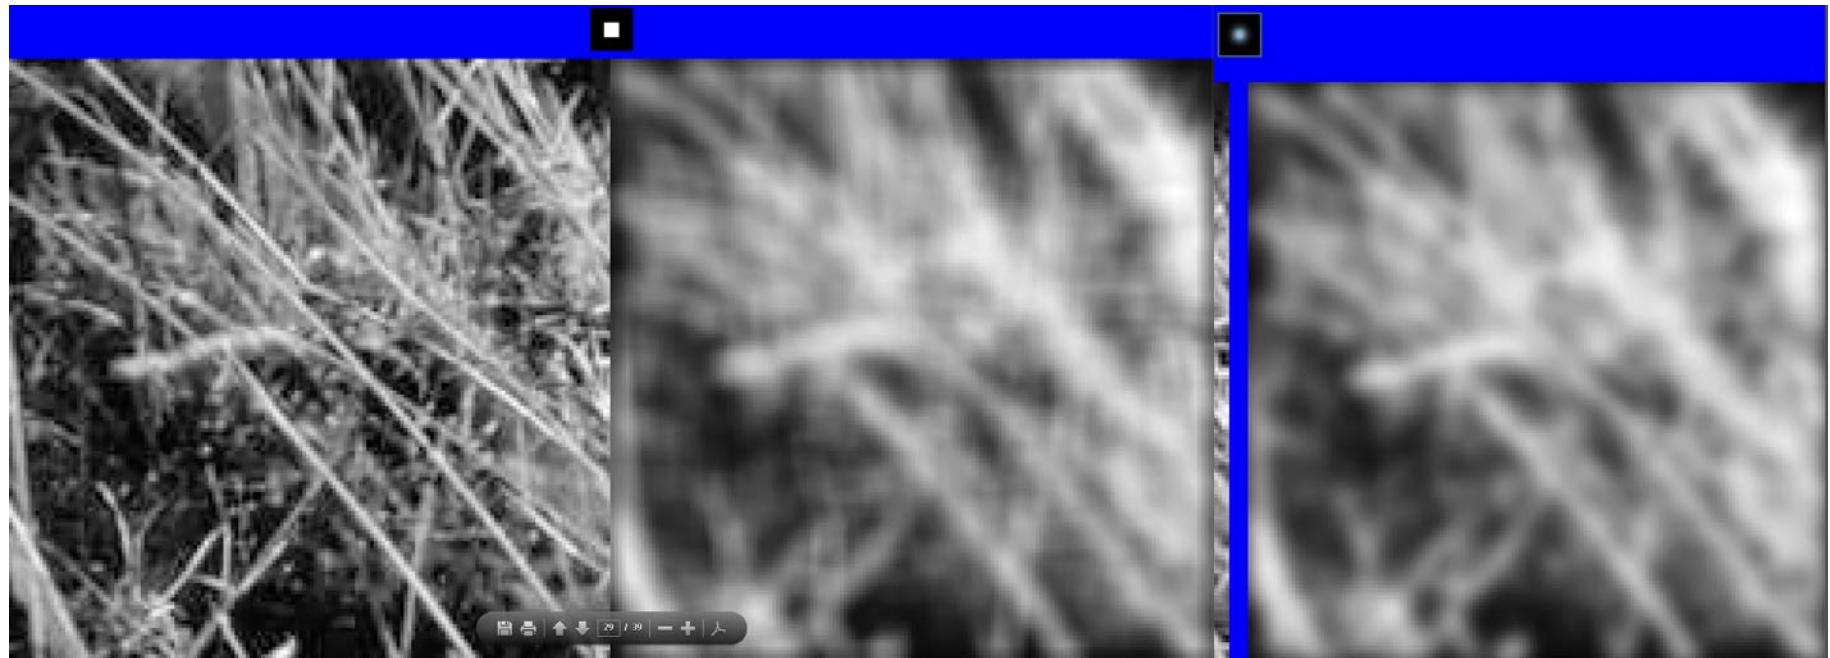


Figure 3.2 The Gaussian function at scales $\sigma = .3$, $\sigma = 1$ and $\sigma = 2$. The kernel is normalized, so the total area under the curve is always unity.

Box versus Gaussian



Convolution

Java demo: <http://www.jhu.edu/signals/convolve/>
<http://www.jhu.edu/signals/discreteconv2/index.html>

- Discrete

$$g(x, y) = w(x, y) * f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

- Continuous

$$g(x, y) = w(x, y) * f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} w(s, t) f(x - s, y - t) ds dt$$

- Same as cross correlation with kernel transposed around each axis
- The two operations (correlation and convolution) are the same if the kernel is symmetric about axes

$$g = w \circ f = w^* * f \quad w^* \text{ reflection of } w$$

Cross-correlation and Convolution

- Cross-correlation

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

- Convolution

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

Convolution: Properties

- Shift invariant, linear
- Commutative

$$f * g = g * f$$

- Associative

$$f * (g * h) = (f * g) * h$$

- Others (discussed later):
 - Derivatives, convolution theorem, spectrum...

Computing Convolution

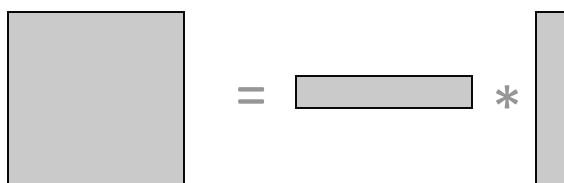
- Compute time
 - $M \times M$ mask
 - $N \times N$ image
- $O(M^2N^2)$ “for” loops are nested 4 deep

Computing Convolution

- Compute time
 - $M \times M$ mask
 - $N \times N$ image
- Special case: *separable*

Two 1D kernels

$$w = \overbrace{w_x * w_y}^{\text{Two 1D kernels}}$$



$$w * f = \underbrace{(w_x * w_y) * f}_{O(M^2N^2)} = \underbrace{w_x * (w_y * f)}_{O(MN^2)}$$

Separable Kernels

- Examples
 - Box/rectangle
 - Bilinear interpolation
 - Combinations of partial derivatives
 - d^2f/dx^2
 - Gaussian
 - Only filter that is both circularly symmetric and separable
- Counter examples
 - Disk
 - Cone
 - Pyramid

Separability

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

```
DisplayTogetherArray[{Plot[gauss[x, σ = 1], {x, -3, 3}],
Plot3D[gauss[x, σ = 1] gauss[y, σ = 1], {x, -3, 3}, {y, -3, 3}]},
ImageSize -> 440];
```

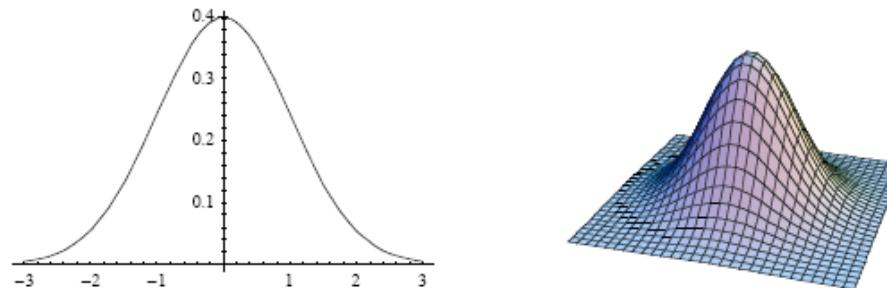
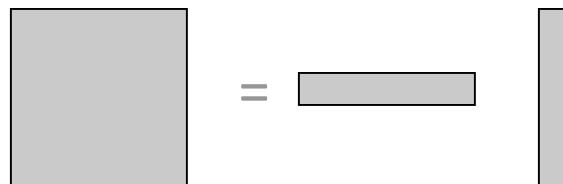


Figure 3.7 A product of Gaussian functions gives a higher dimensional Gaussian function. This is a consequence of the separability.

Key Concepts

- Separability

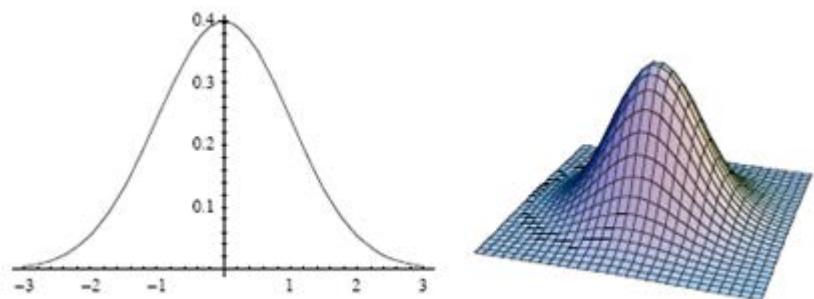
$$w = w_x * w_y$$



$$w * f = (w_x * w_y) * f = w_x * (w_y * f)$$

$$\underbrace{\hspace{10em}}_{O(M^2N^2)} \quad \underbrace{\hspace{10em}}_{O(MN^2)}$$

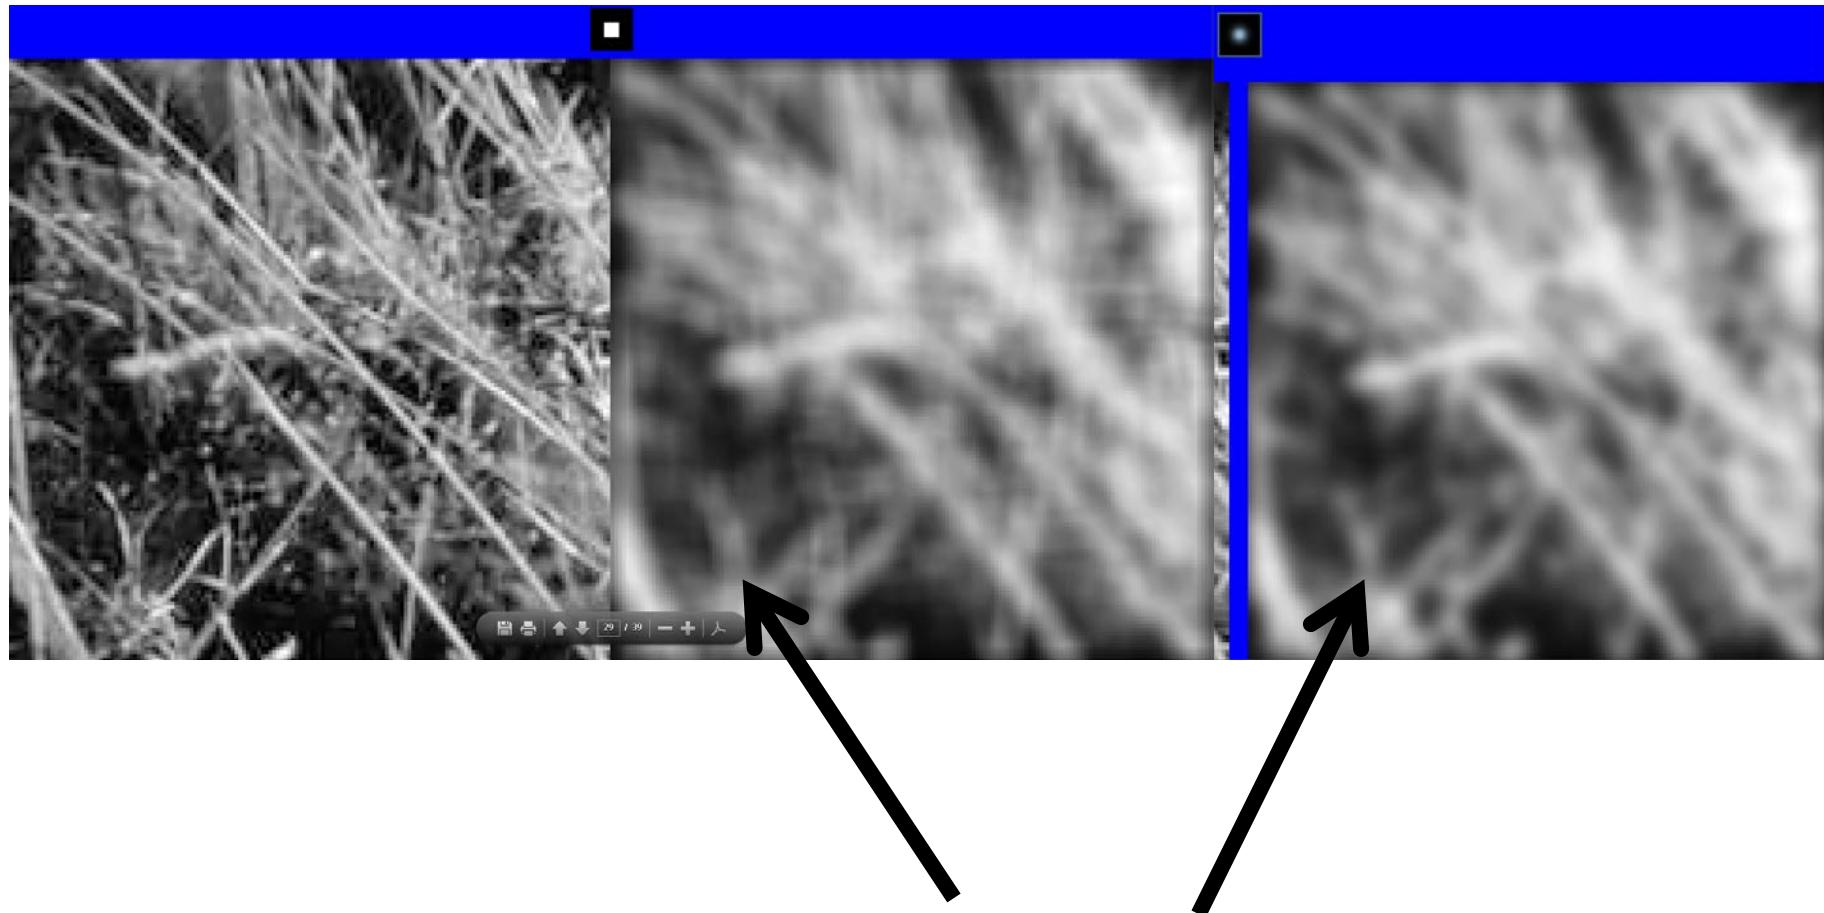
$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$



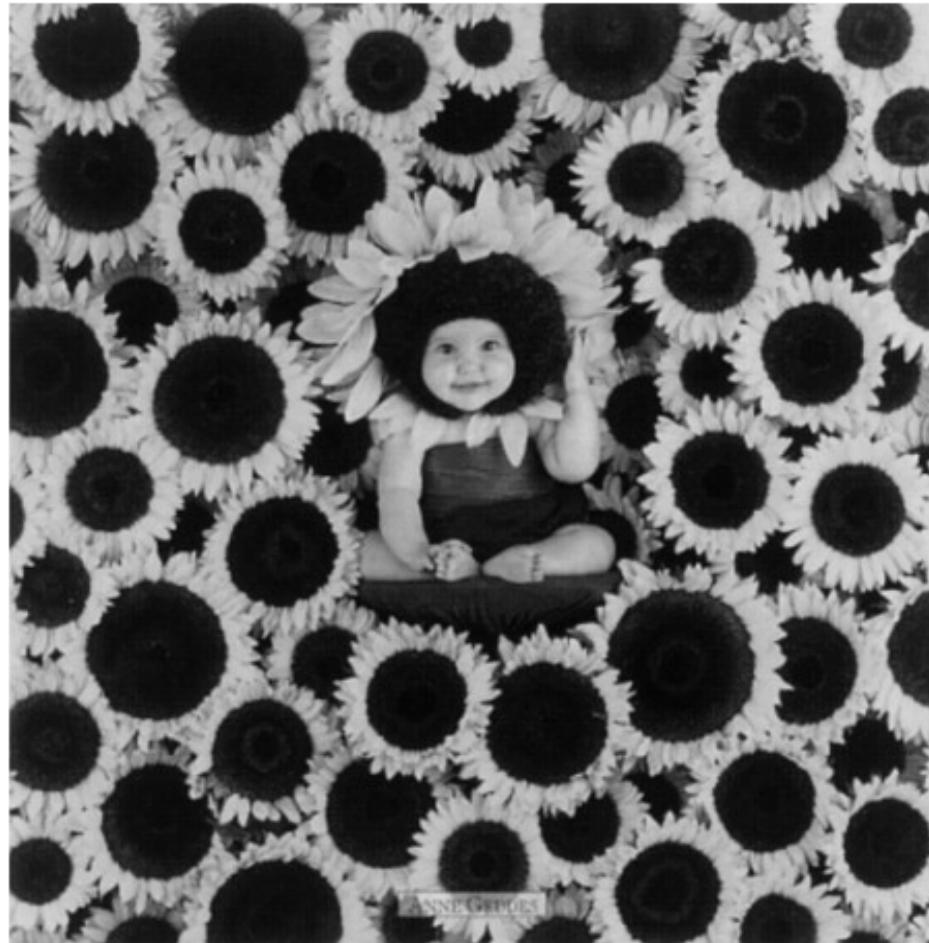
The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

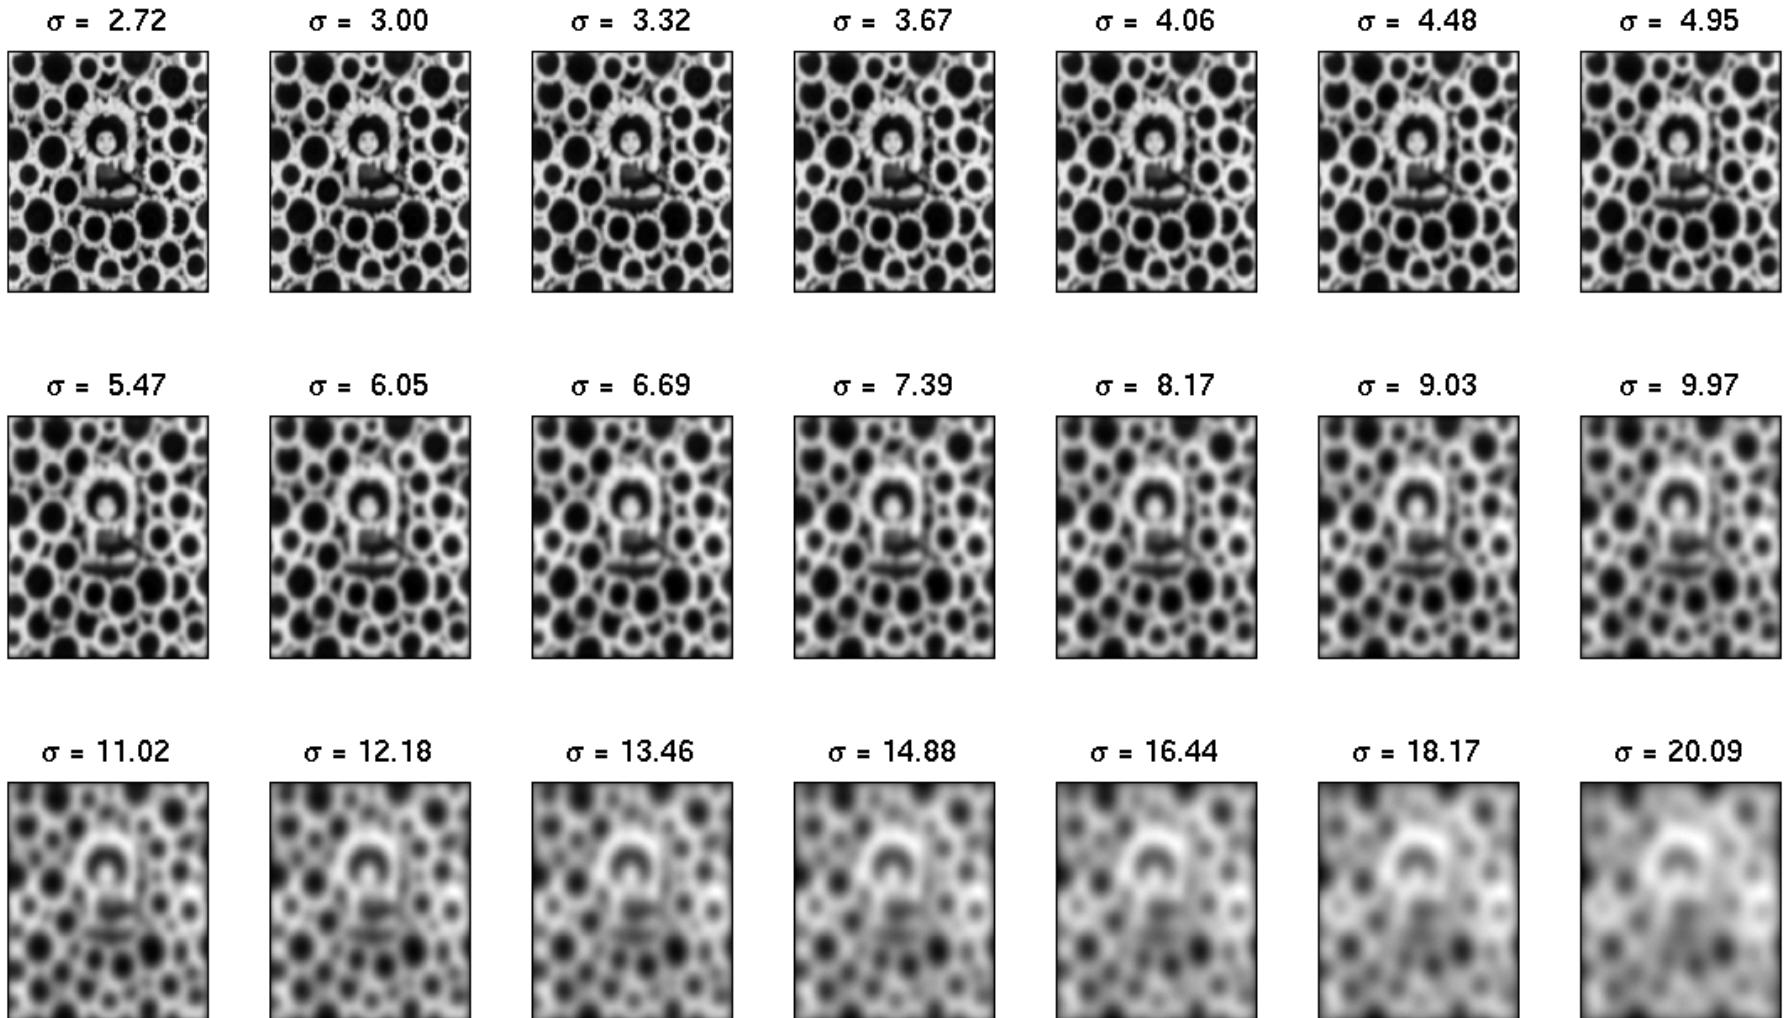
Box versus Gaussian



Gaussian Filtering



Gaussian Filtering

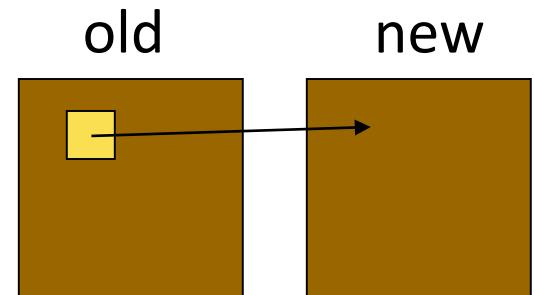


Nonlinear Methods For Filtering

- Median filtering
- Bilateral filtering
- Neighborhood statistics and nonlocal filtering

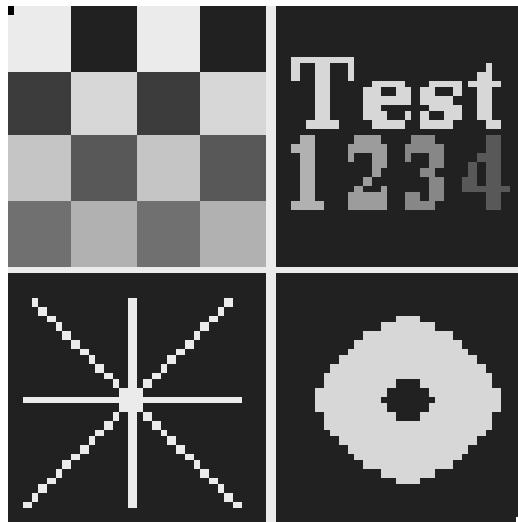
Median Filtering

- For each neighborhood in image
 - Sliding window
 - Usually odd size (symmetric) 5x5, 7x7,...
- Sort the greyscale values
- Set the center pixel to the median
- Important: use “Jacobi” updates
 - Separate input and output buffers
 - All statistics on the original image

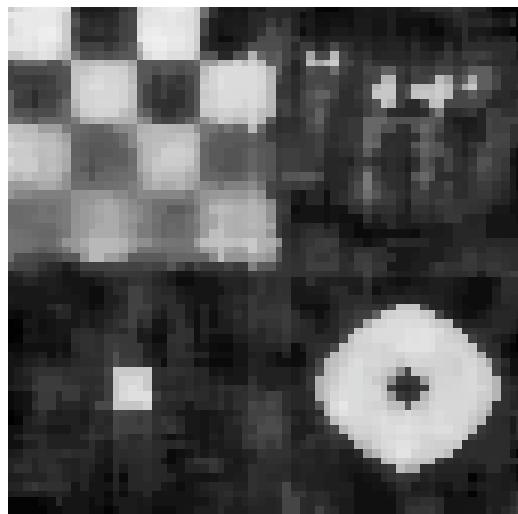


Median vs Gaussian

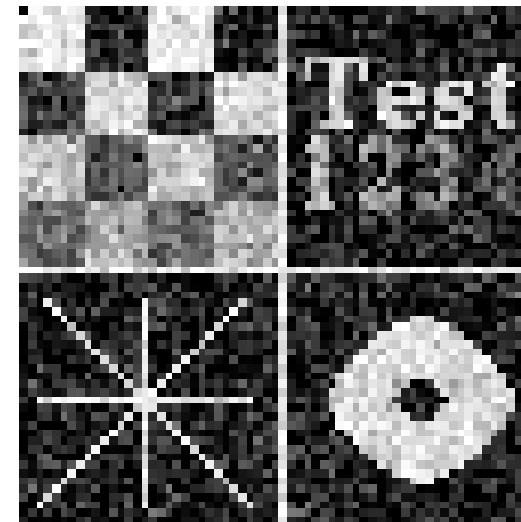
Original



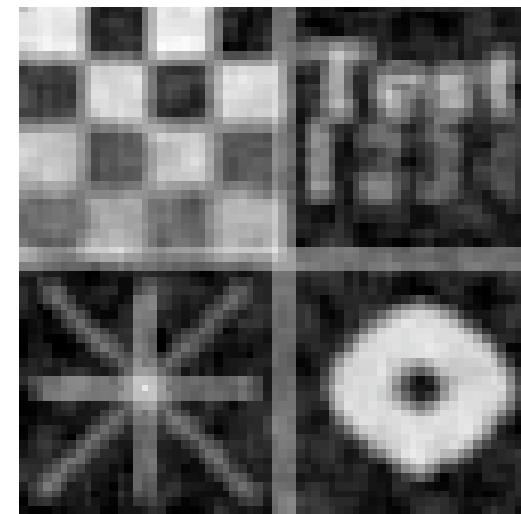
3x3 Median



+ Gaussian
Noise



3x3 Box



Median Filter

- **Issues**
 - Boundaries
 - Compute on pixels that fall within window
 - Computational efficiency
 - What is the best algorithm?
- **Properties**
 - Removes outliers (replacement noise – salt and pepper)
 - Window size controls size of structures
 - Preserves straight edges, but rounds corners and features

Replacement Noise

- Also: “shot noise”, “salt&pepper”
- Replace certain % of pixels with samples from pdf
- Best filtering strategy: filter to avoid outliers



Smoothing of S&P Noise

- It's not zero mean (locally)
- Averaging produces local biases



Smoothing of S&P Noise

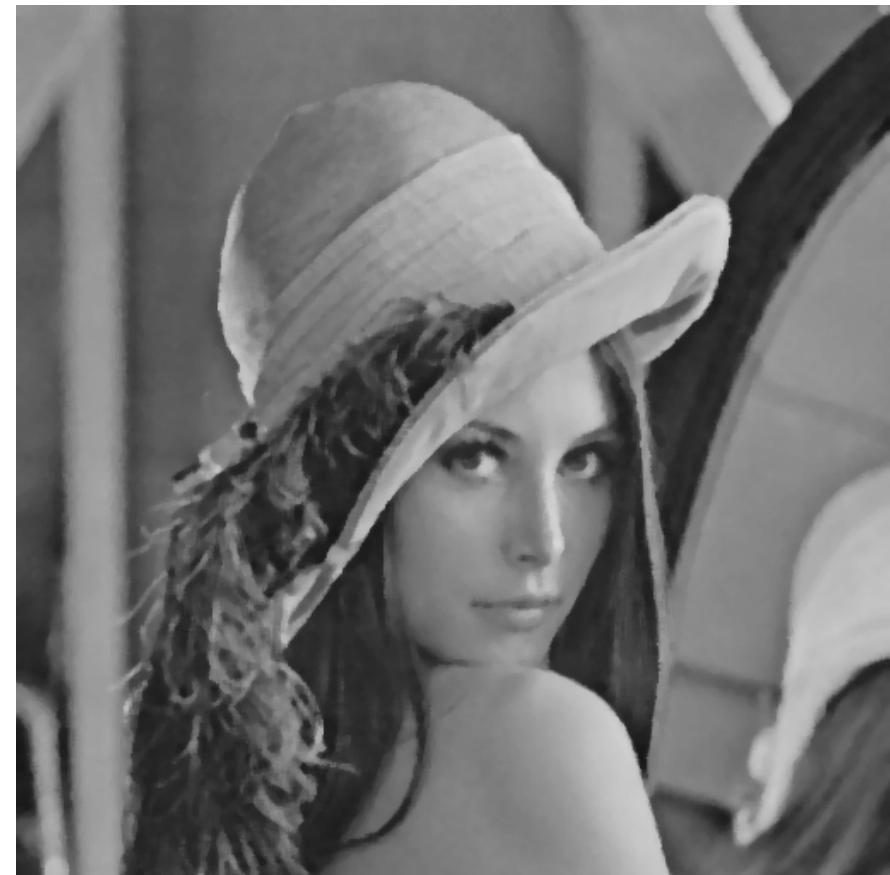
- It's not zero mean (locally)
- Averaging produces local biases



Median Filtering



Median 3x3



Median 5x5

Median Filtering



Median 3x3



Median 5x5

Median Filtering

- Iterate



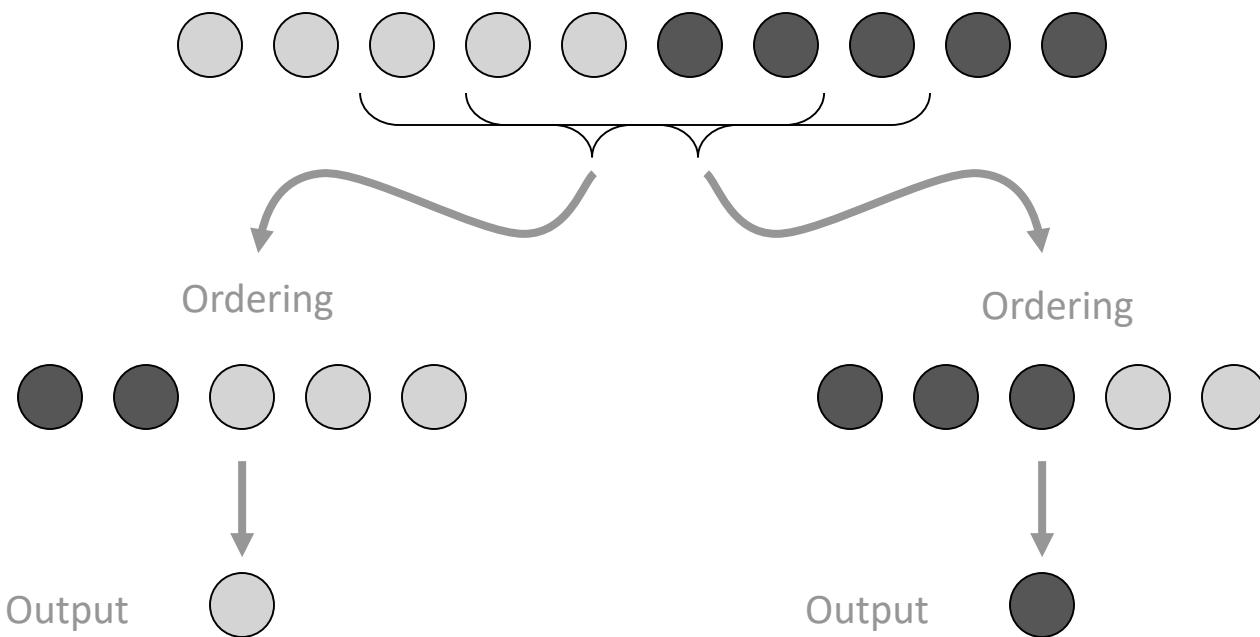
Median 3x3



2x Median 3x3

Median Filtering

- Image model: piecewise constant (flat)



Order Statistics

- Median is special case of order-statistics filters
- Instead of weights based on neighborhoods, weights are based on ordering of data

Neighborhood	Ordering
X_1, X_2, \dots, X_N	$X_{(1)} \leq X_{(2)} \leq \dots \leq X_{(N)}$

Filter $F(X_1, X_2, \dots, X_N) = \alpha_1 X_{(1)} + \alpha_2 X_{(2)} + \dots + \alpha_N X_{(N)}$

Neighborhood average (box)

$$\alpha_i = 1/N$$

Median filter

$$\alpha_i = \begin{cases} 1 & i = (N + 1)/2 \\ 0 & \text{otherwise} \end{cases}$$

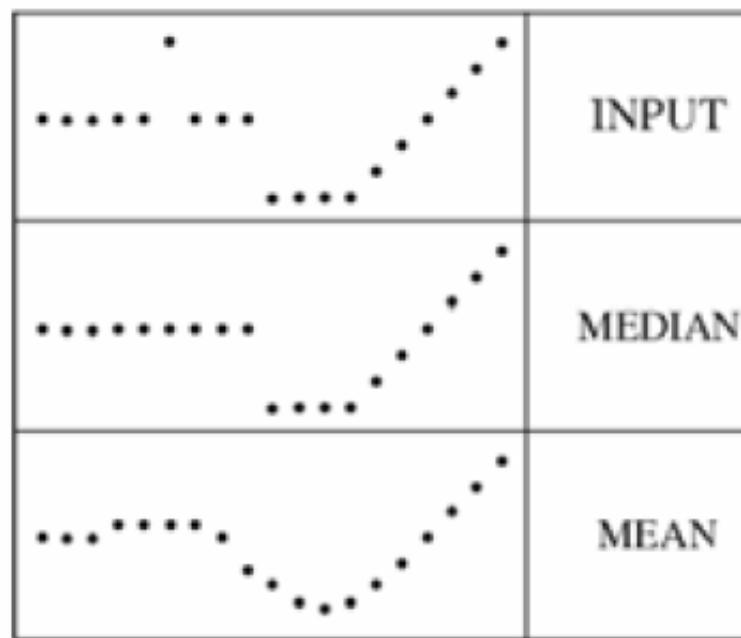
Trimmed average (outlier removal)

$$\alpha_i = \begin{cases} 1/M & (N - M + 1)/2 \leq i \leq (N + M + 1)/2 \\ 0 & \text{otherwise} \end{cases}$$

Median filter

- What advantage does median filtering have over Gaussian filtering?
 - Robustness to outliers

filters have width 5 :



Source: K. Grauman

Piecewise Flat Image Models

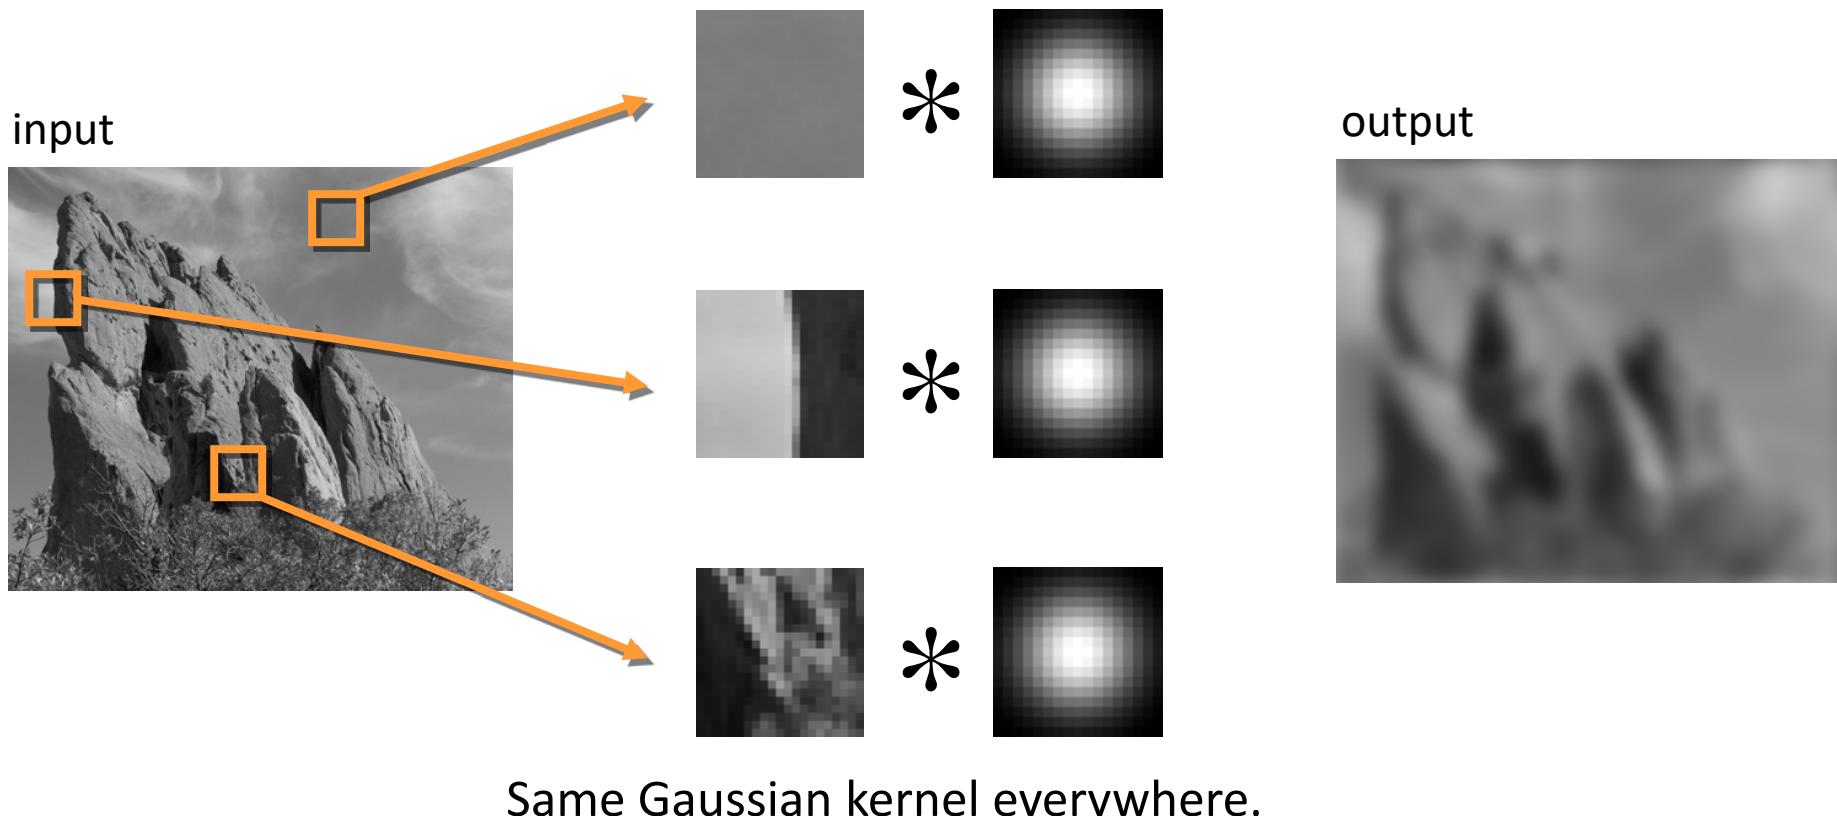
- Image piecewise flat -> average only within similar regions
- Problem: don't know region boundaries



Piecewise-Flat Image Models

- Assign probabilities to other pixels in the image belonging to the same region
- Two considerations
 - Distance: far away pixels are less likely to be same region
 - Intensity: pixels with different intensities are less likely to be same region

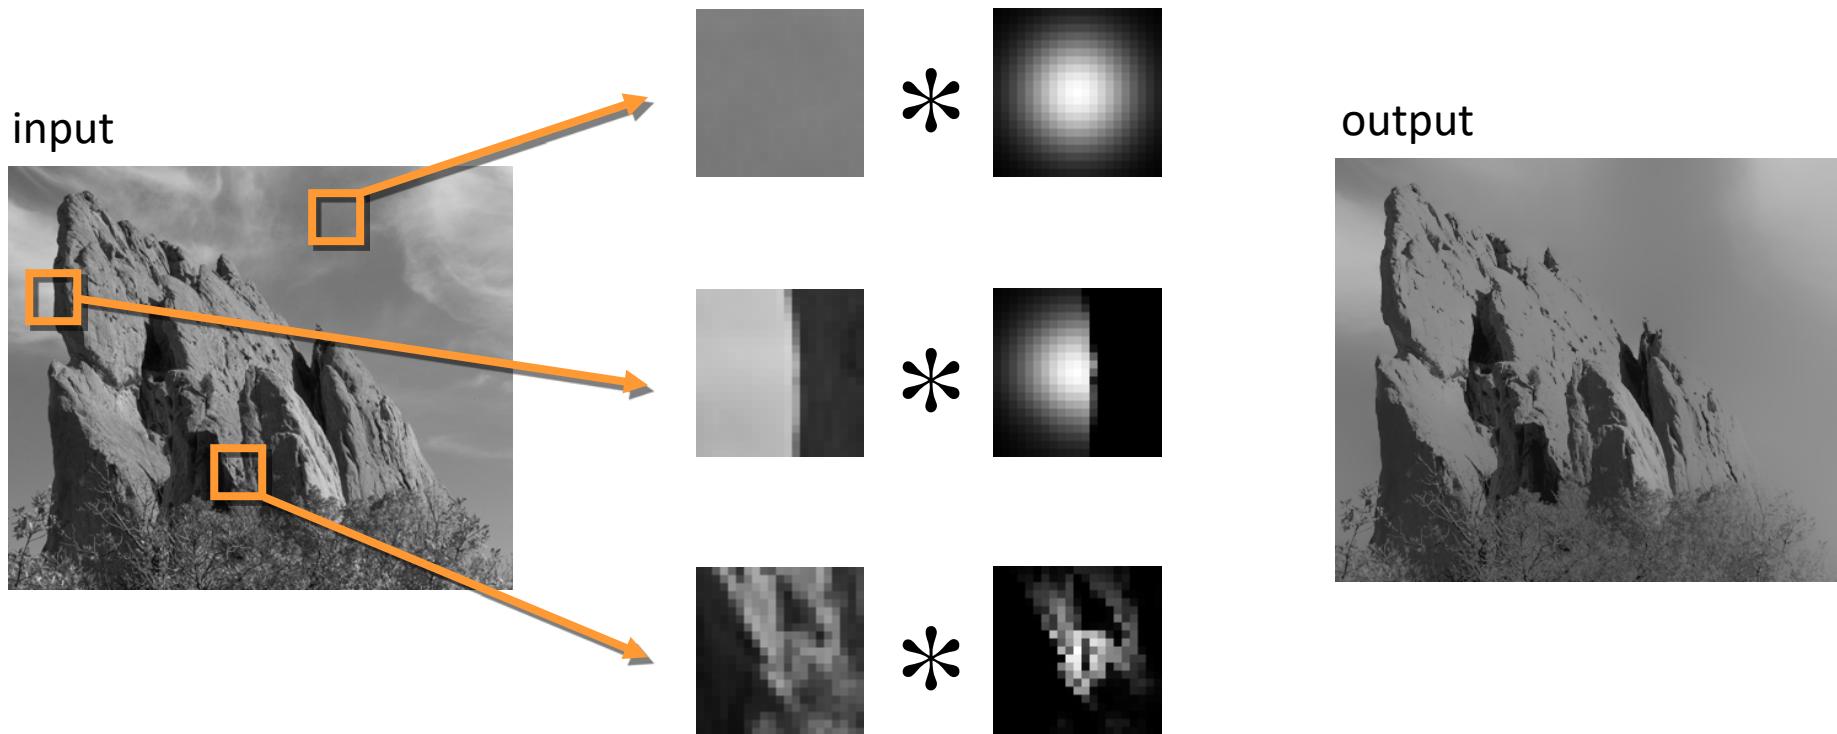
Gaussian: Blur Comes from Averaging across Edges



Bilateral Filter

No Averaging across Edges

[Aurich 95, Smith 97, Tomasi 98]



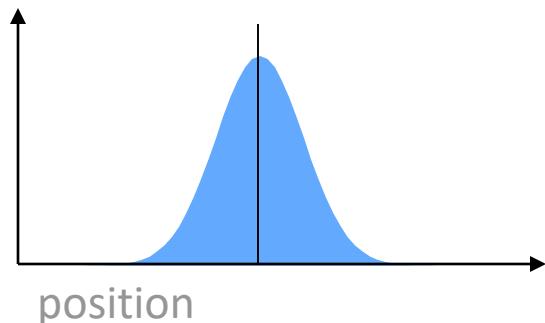
The kernel shape depends on the image content.

Main Idea

Distance (kernel/pdf)

$$G(\mathbf{x}_i - \mathbf{x}_j)$$

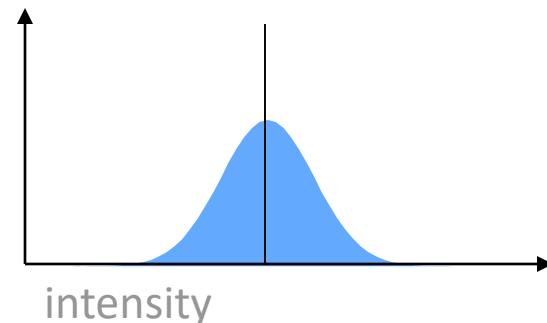
Prob
pixel
belongs
to same
region as
i



Distance (pdf)

$$H(f_i - f_j)$$

Prob
pixel
belongs
to same
region as
i



Bilateral Filter Definition: an Additional Edge Term

Same idea: **weighted average of pixels.**

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s} (|| p - q ||) G_{\sigma_r} (| I_p - I_q |) I_q$$

new
not new
new

normalization factor
space weight
range weight

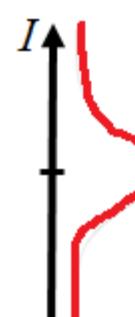
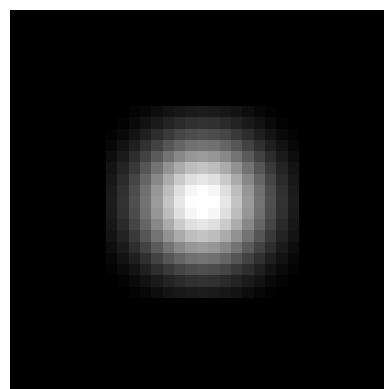
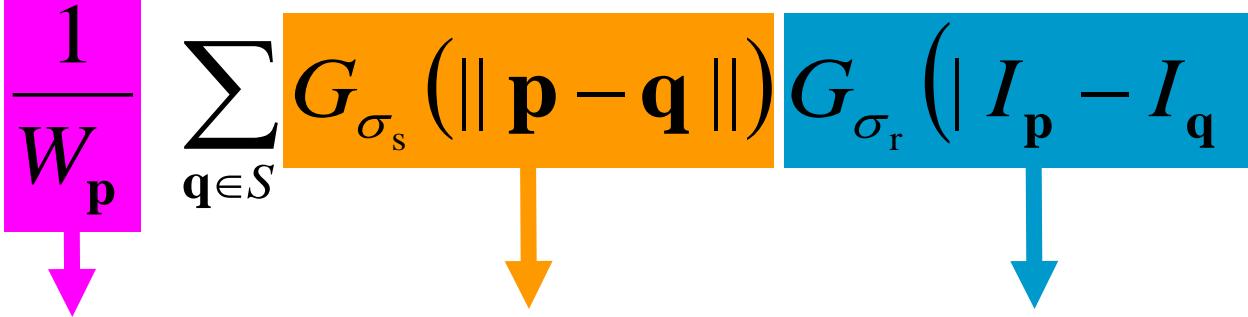
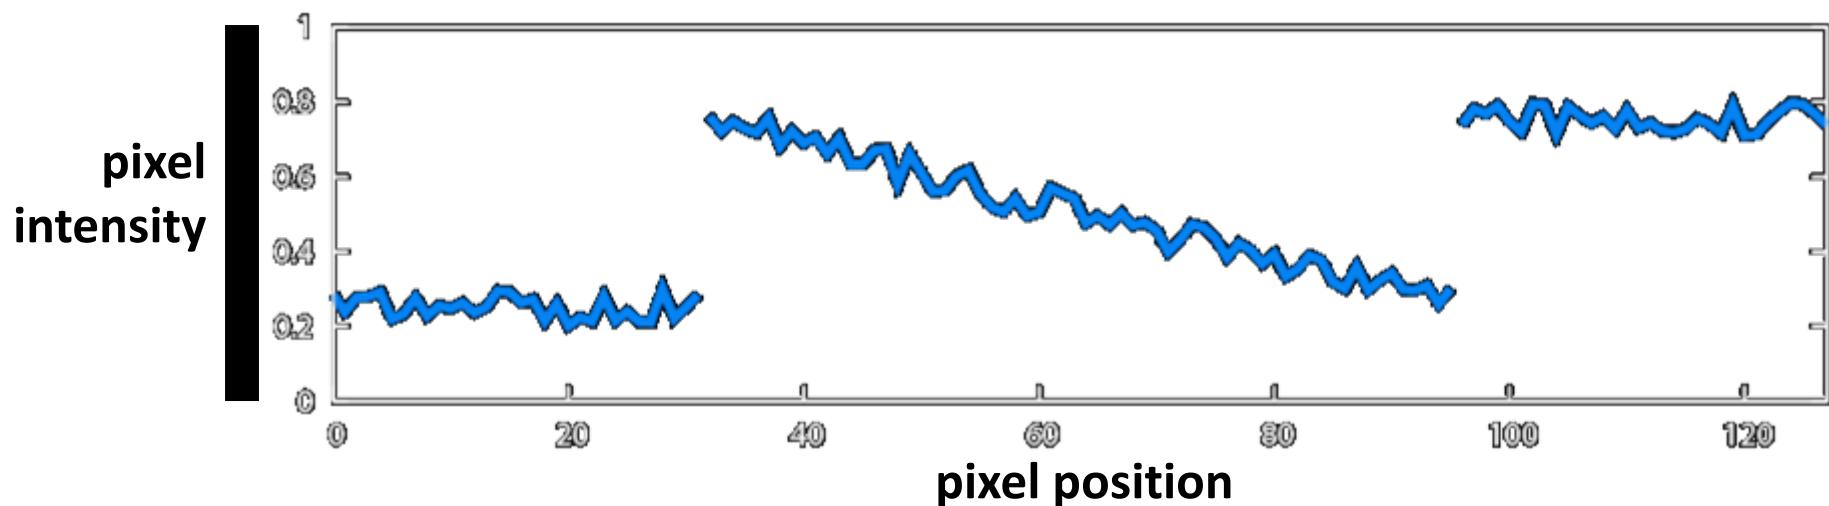


Illustration a 1D Image

- 1D image = line of pixels

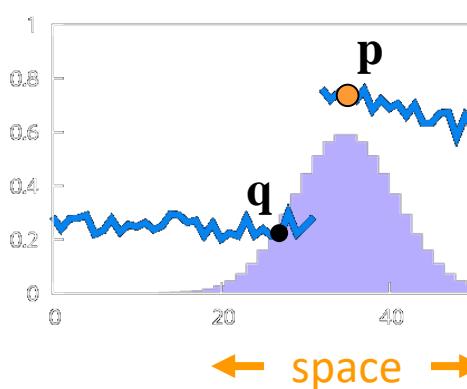


- Better visualized as a plot



Gaussian Blur and Bilateral Filter

Gaussian blur

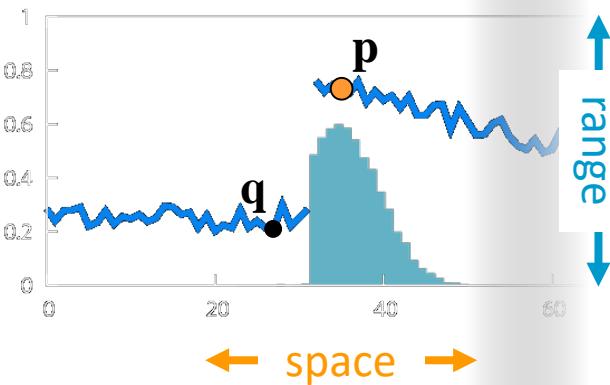


$$GB[I]_p = \sum_{q \in S} G_\sigma(\|p - q\|) I_q$$

space

Bilateral filter

[Aurich 95, Smith 97, Tomasi 98]



$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

space range

normalization

Bilateral Filter

- Neighborhood – sliding window
- Weight contribution of neighbors according to:

$$f_i \leftarrow k_i^{-1} \sum_{j \in N} f_j G(\mathbf{x}_i - \mathbf{x}_j) H(f_i - f_j)$$

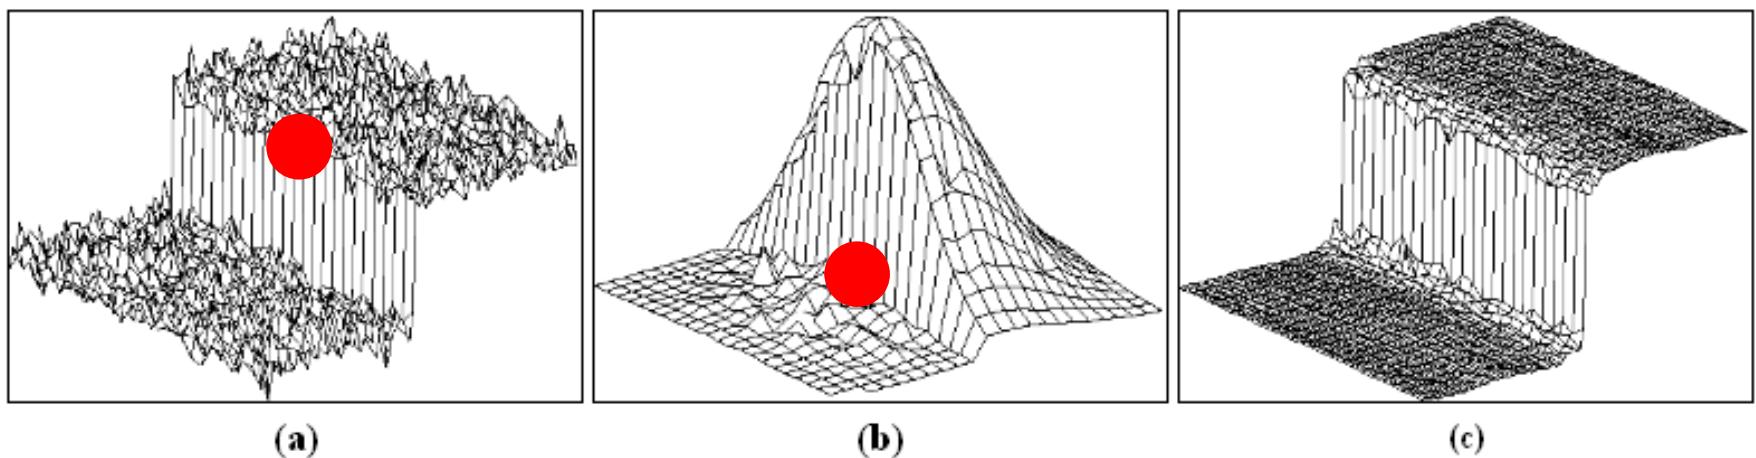
$$k_i = \sum_{j \in N} G(\mathbf{x}_i - \mathbf{x}_j) H(f_i - f_j)$$

normalization: all weights add up to 1

- G is a Gaussian (or lowpass), as is H , N is neighborhood,
 - Often use $G(r_{ij})$ where r_{ij} is distance between pixels
 - Update must be normalized for the samples used in this (particular) summation
- Spatial Gaussian with extra weighting for intensity
 - Weighted average in neighborhood with downgrading of intensity outliers

Bilateral Filter

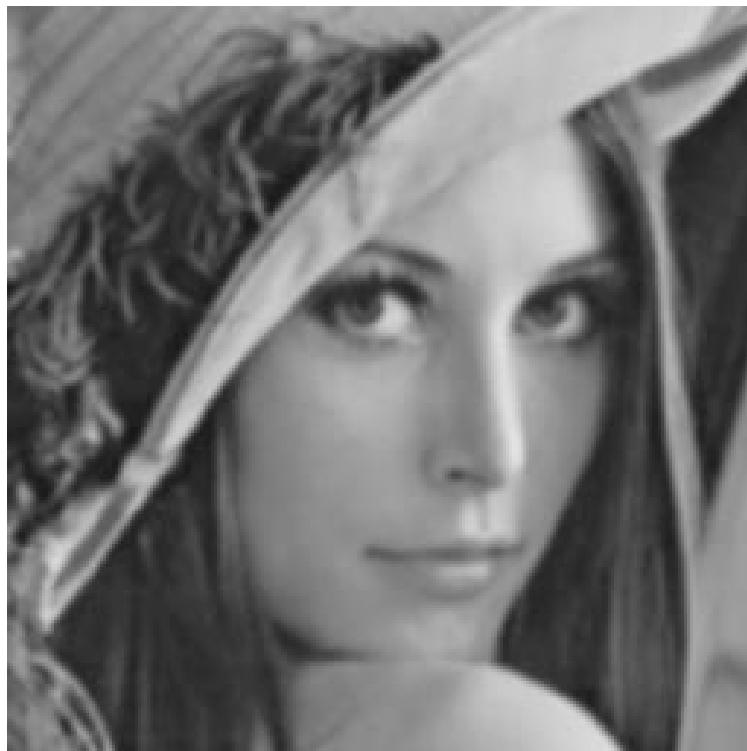
Replaces the pixel value at x with an average of similar and nearby pixel values.



When the bilateral filter is centered, say, on a pixel on the bright side of the boundary, the similarity function s assumes values close to one for pixels on the same side, and values close to zero for pixels on the dark side. The similarity function is shown in figure 1(b) for a 23x23 filter support centered two pixels to the right of the step in figure 1(a).

Bilateral Filtering

Replaces the pixel value at x with an average of similar and nearby pixel values.



Gaussian Blurring



Bilateral

Bilateral Filtering



Gaussian Blurring

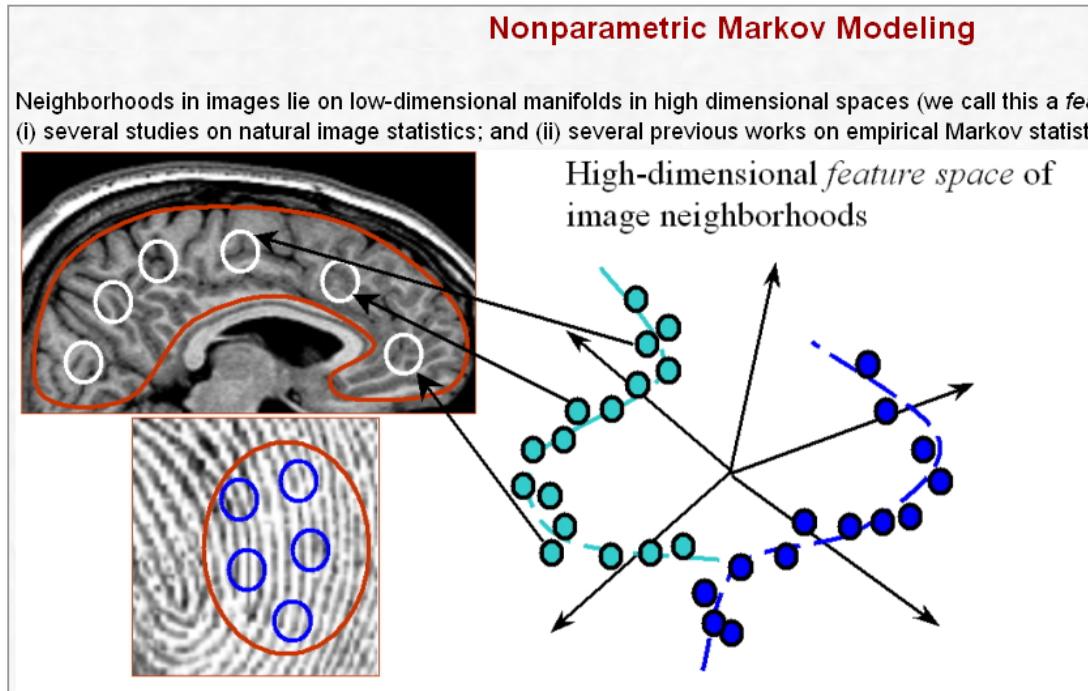


Bilateral

Nonlocal Averaging

- Recent algorithm
 - NL-means, Baudes et al., 2005
 - UNTA, Awate & Whitaker, 2005
- Different model
 - No need for piecewise-flat
 - Images consist of some set of pixels with similar neighborhoods → average several of those
 - Scattered around
 - General area of a pixel
 - All around
- Idea
 - Average sets of pixels with similar neighborhoods

UINTA: Unsupervised Information-Theoretic Adaptive Filtering : Excellent Introduction and Additional Readings (Suyash P. Awate)



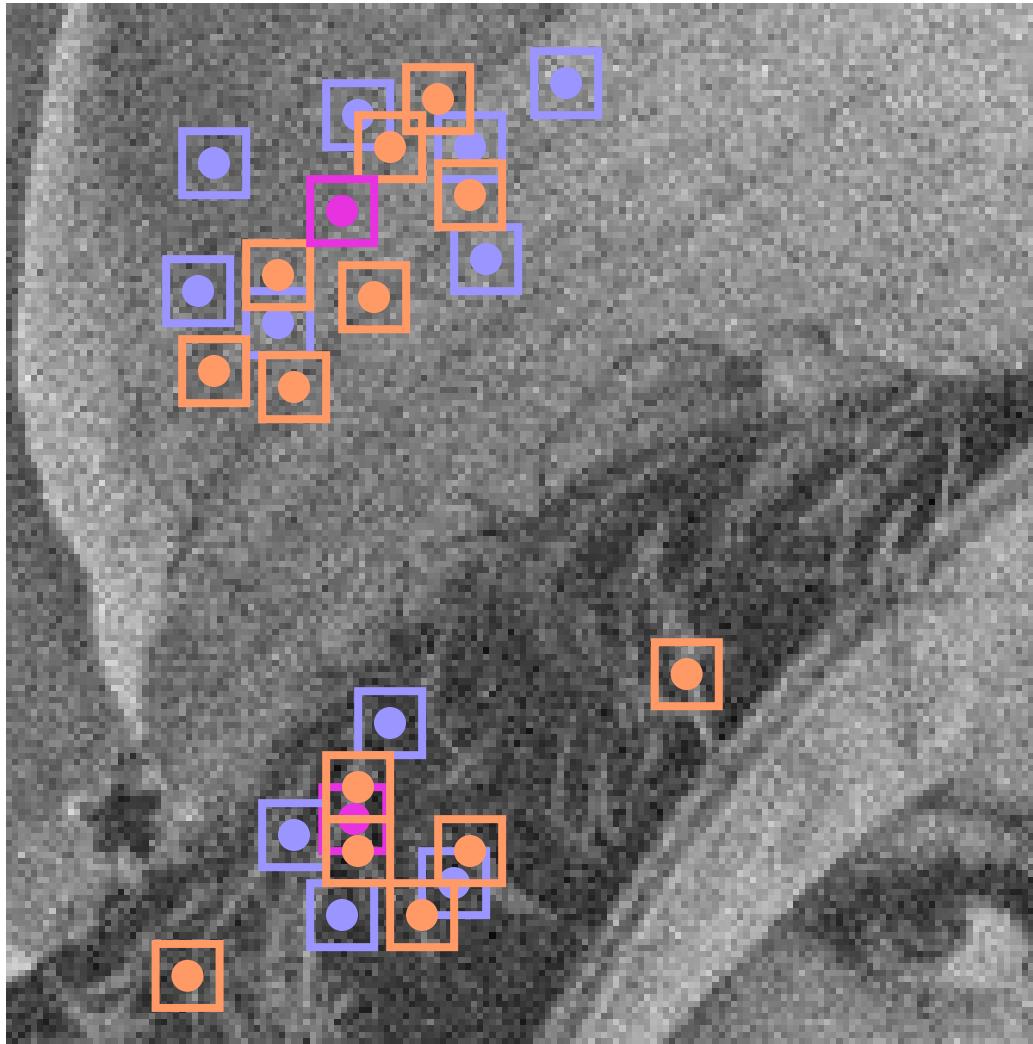
<http://www.cs.utah.edu/~suyash/pubs/uinta/>

Suyash P. Awate, Ross T. Whitaker

Unsupervised, Information-Theoretic, Adaptive Image Filtering with Applications to Image Restoration

IEEE Trans. Pattern Analysis & Machine Intelligence (TPAMI) 2006, Vol. 28, Num. 3, pp. 364-376

Nonlocal Averaging



Some Details

- Window sizes: good range is $5 \times 5 \rightarrow 11 \times 11$
- How to choose samples:
 - Random samples from around the image
 - UNTA, Awate&Whitaker
 - Block around pixel (bigger than window, e.g. 51×51)
 - NL-means
- Iterate
 - UNTA: smaller updates and iterate

NL-Means Algorithm

- For each pixel, p
 - Loop over set of pixels nearby
 - Compare the neighborhoods of those pixels to the neighborhood of p and construct a set of weights
 - Replace the value of p with a weighted combination of values of other pixels
- Repeat... but 1 iteration is pretty good

Results



Noisy image (range 0.0-1.0)



Bilateral filter (3.0, 0.1)

Results



Bilateral filter (3.0, 0.1)



NL means (7, 31, 1.0)

Results



Bilateral filter (3.0, 0.1)



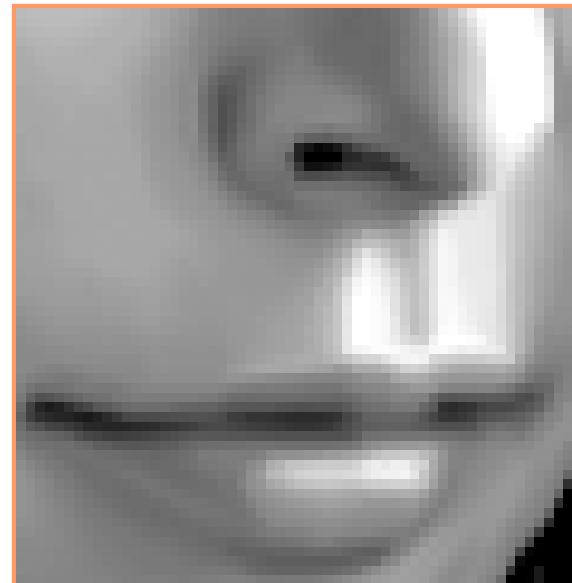
NL means (7, 31, 1.0)

Less Noisy Example

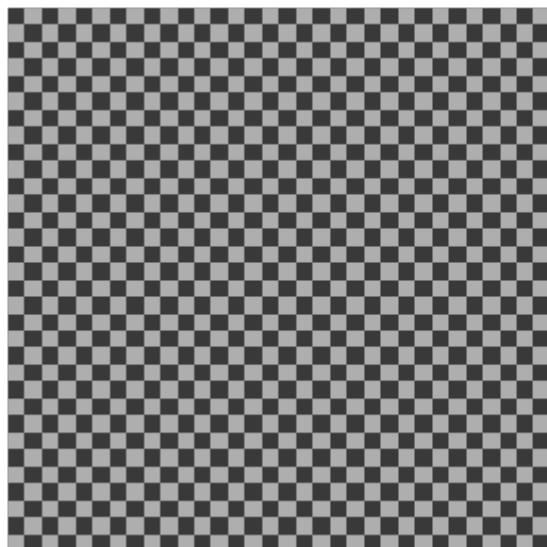




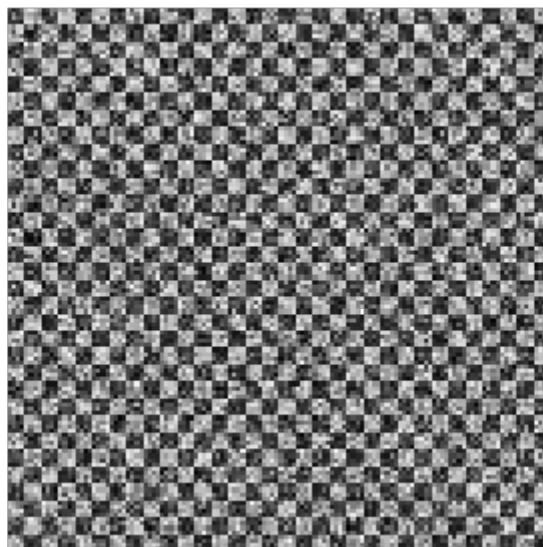
Less Noisy Example



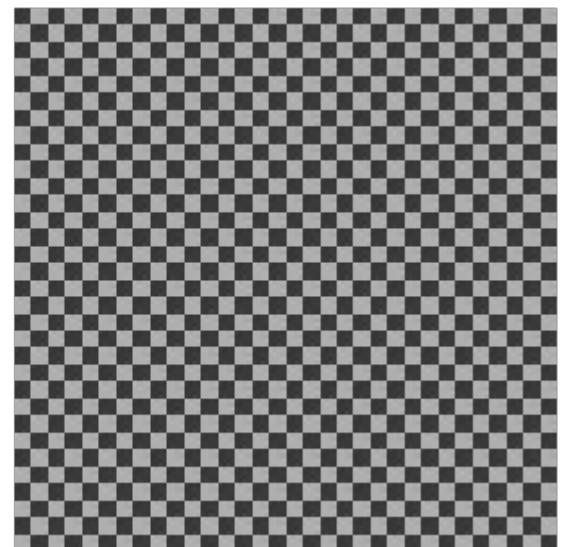
Results



Original

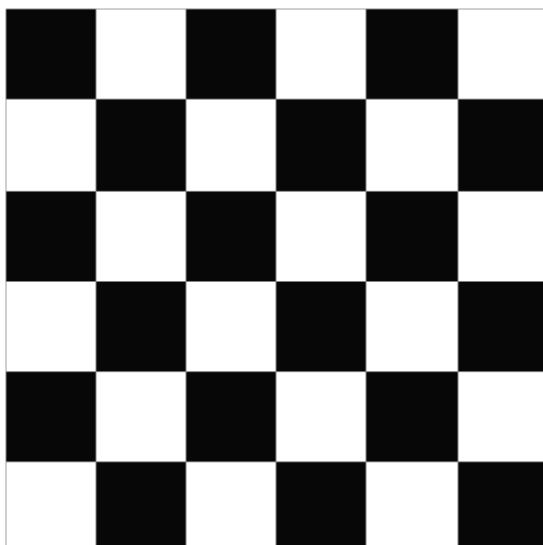


Noisy

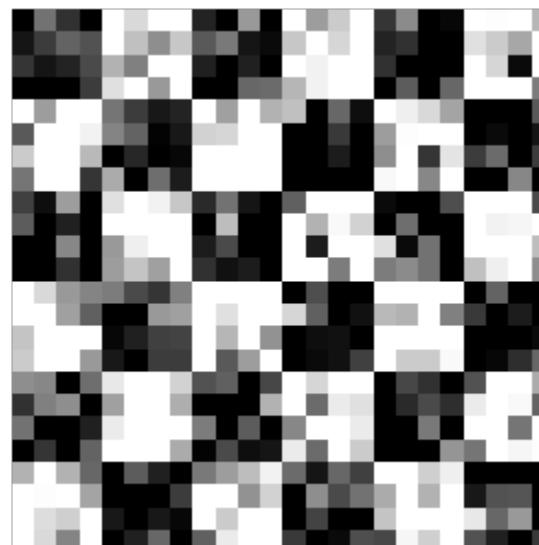


Filtered

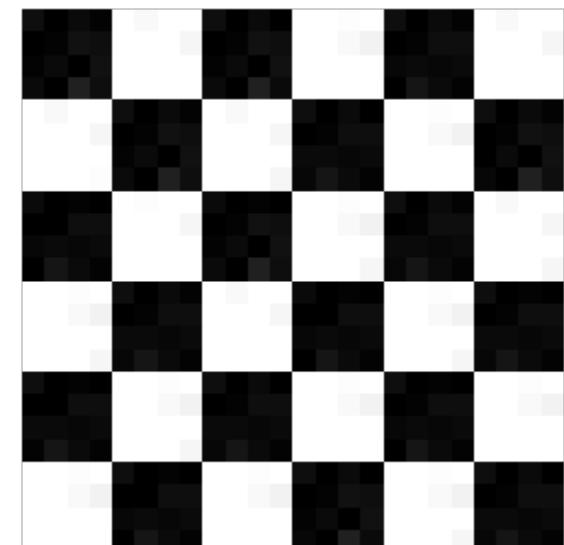
Checkerboard With Noise



Original

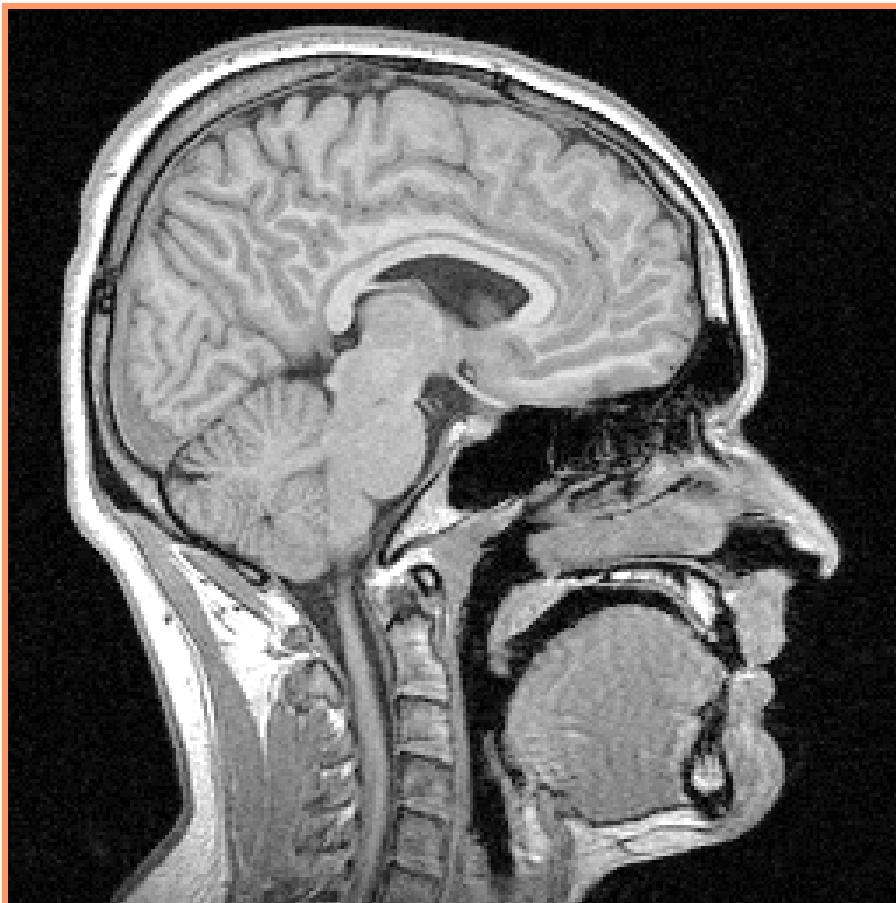


Noisy



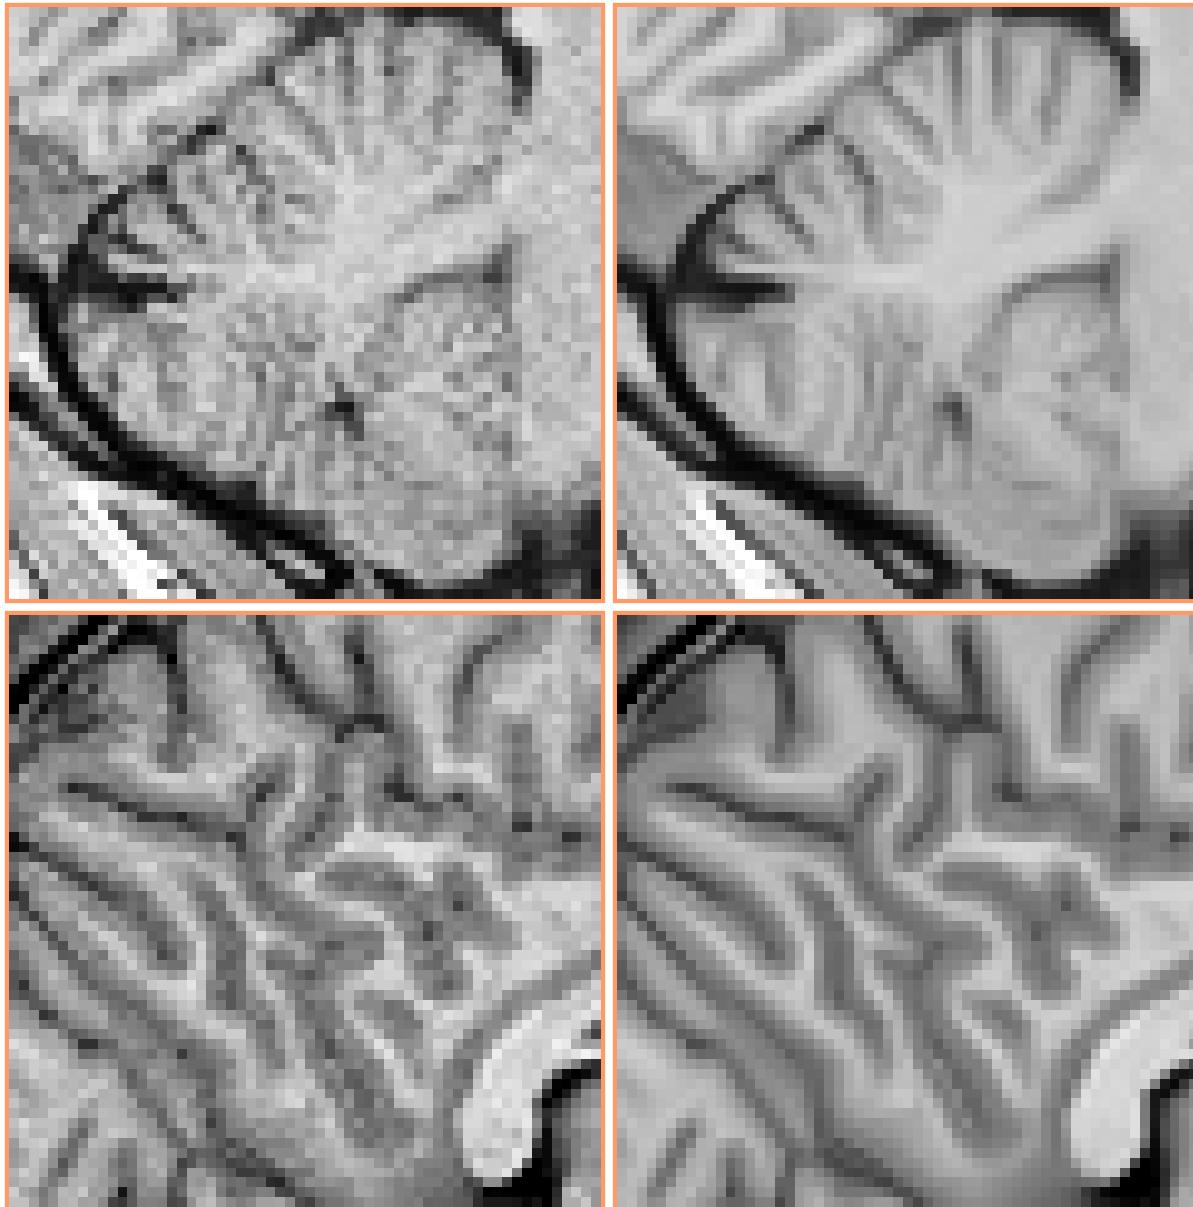
Filtered

MRI Head

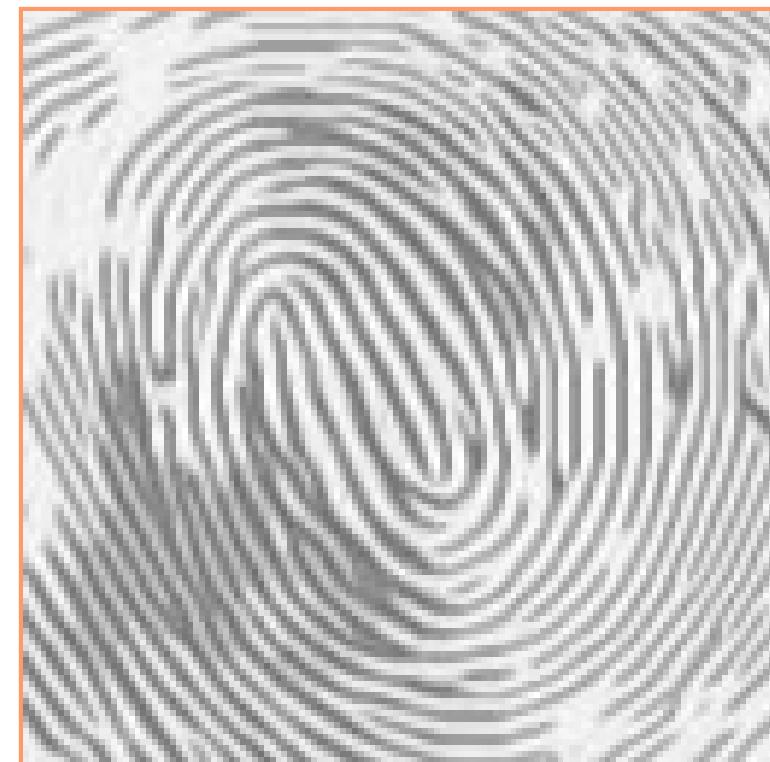




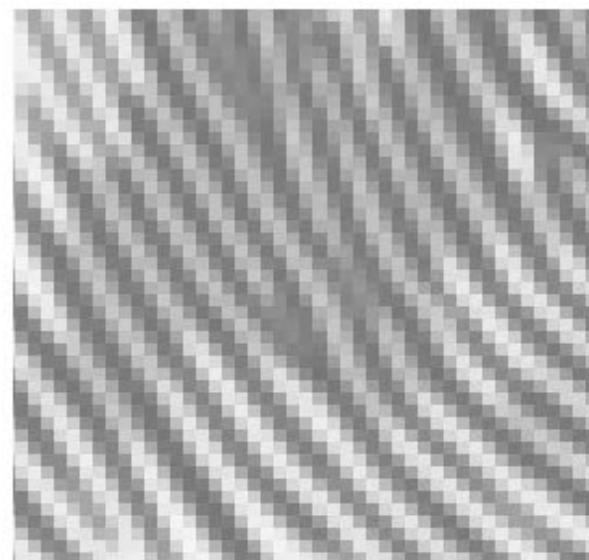
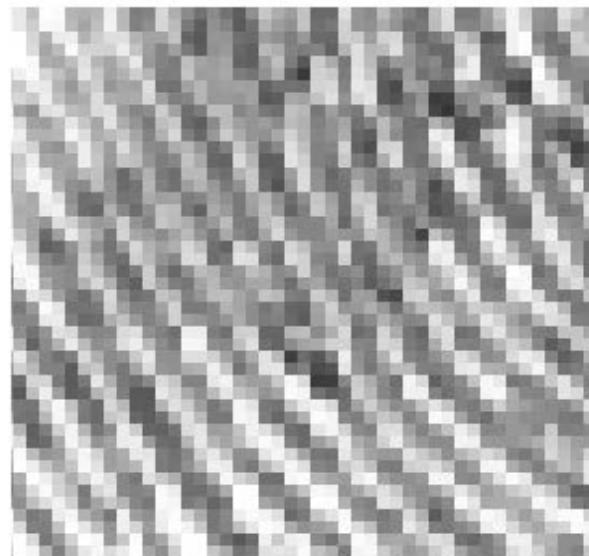
MRI Head



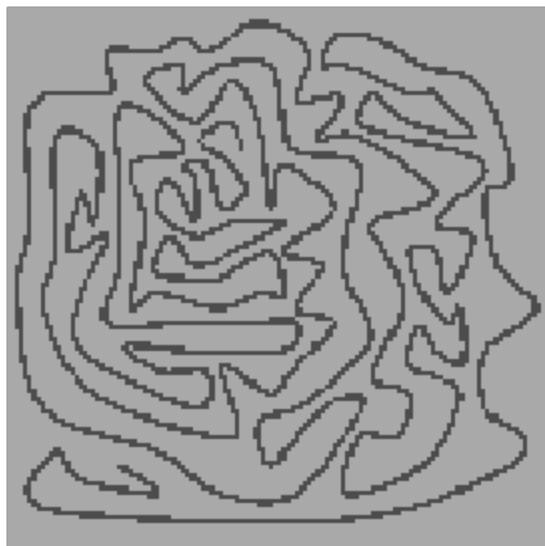
Fingerprint



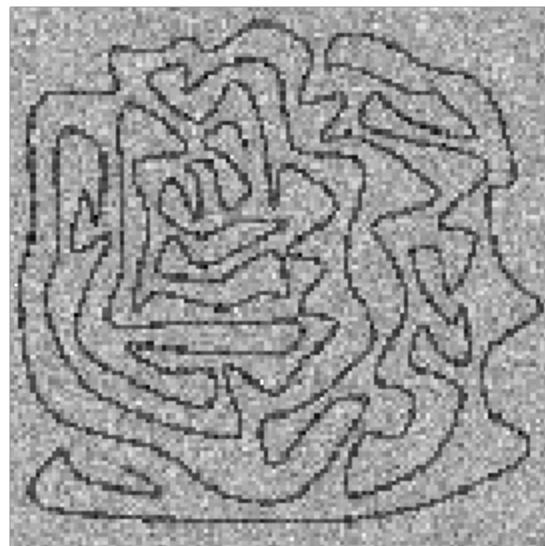
Fingerprint



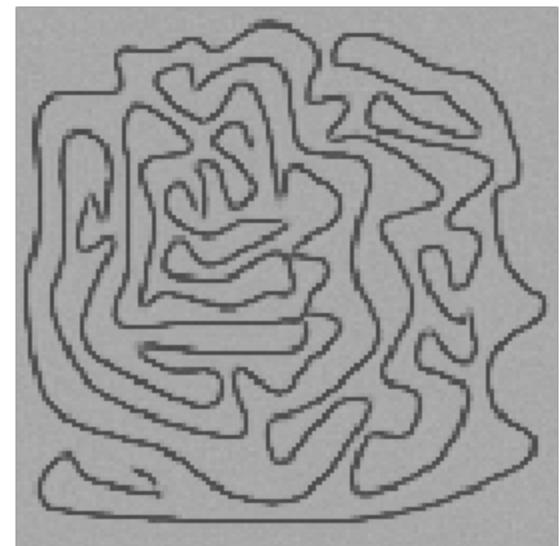
Results



Original



Noisy



Filtered

Results



Original

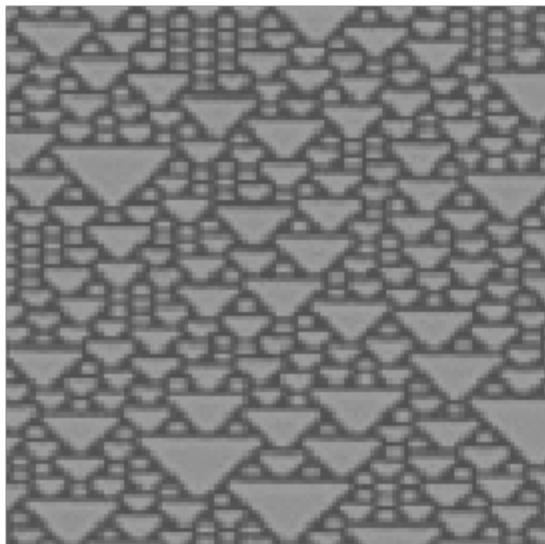


Noisy

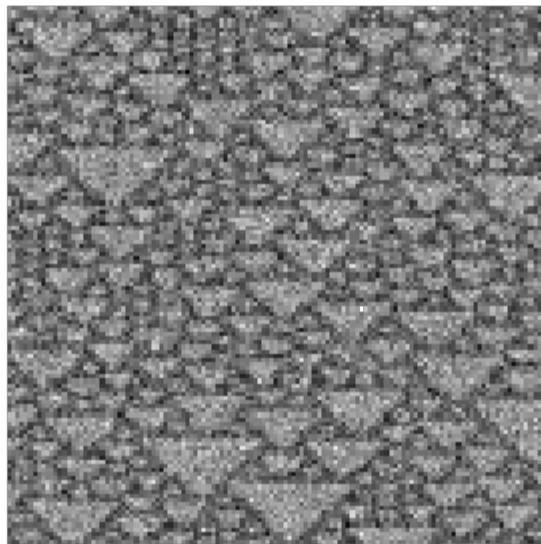


Filtered

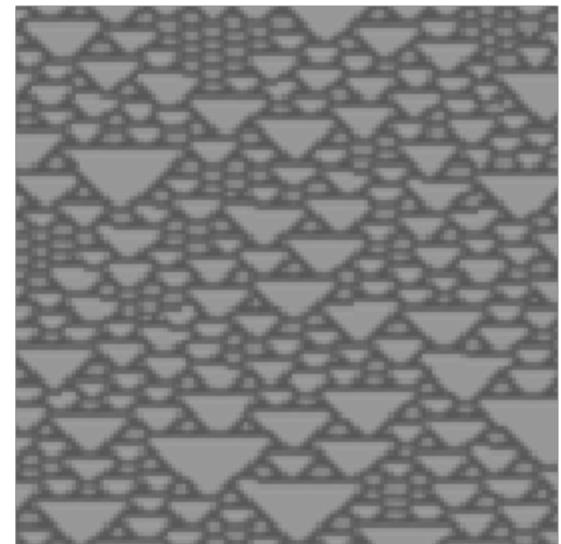
Results



Original

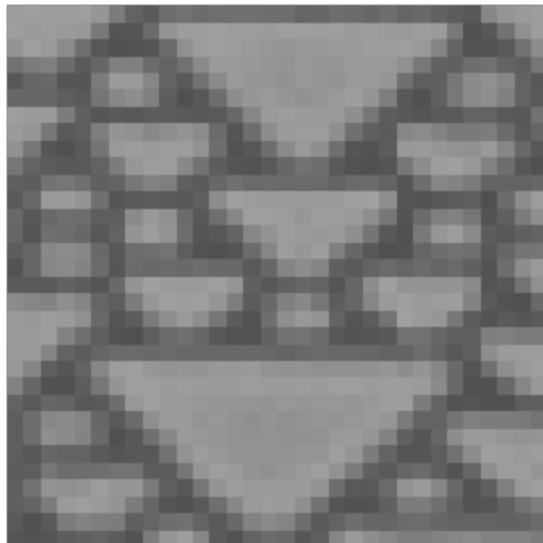


Noisy

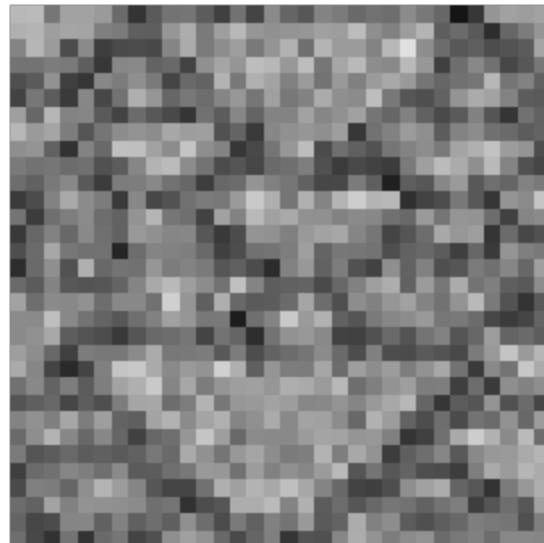


Filtered

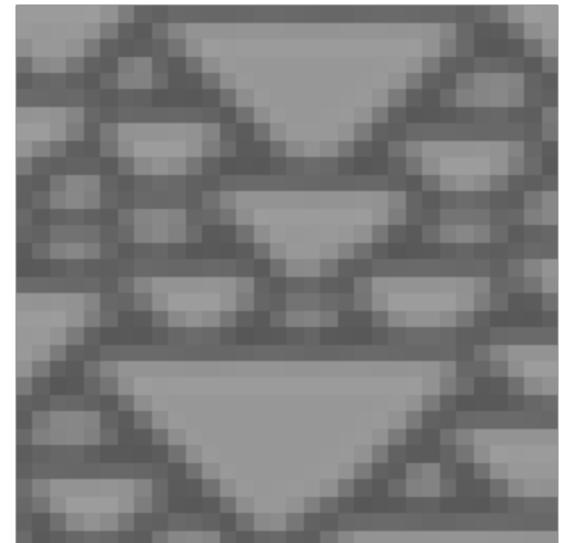
Fractal



Original



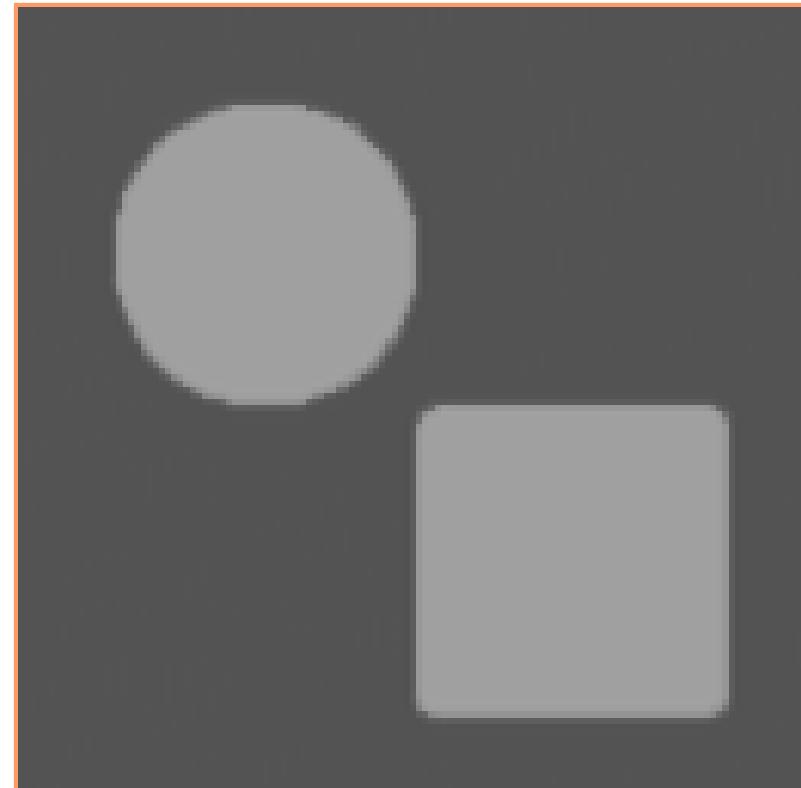
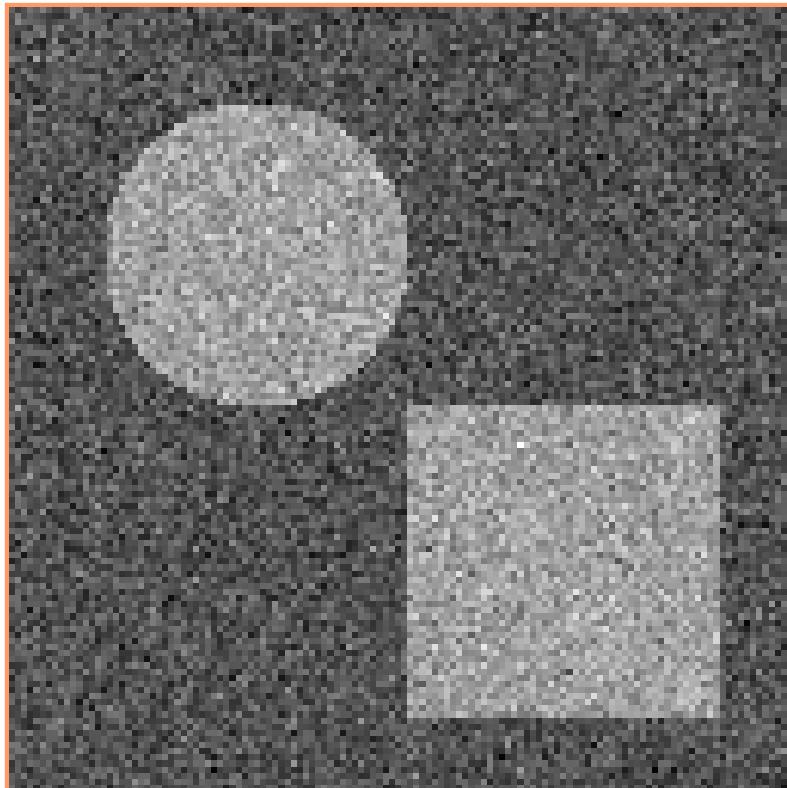
Noisy



Filtered

Piecewise Constant

- Several 10s of Iterations
- Tends to obliterate rare events (e.g. corners)



Texture, Structure

